

WHO ATE MY IOPS? IDENTIFYING HEAVY HITTERS USING JOBSTATS

James Beal, Principal Systems Administrator

Dave Holland, Principal Systems Administrator

V07 2022/11/30

AGENDA



INTRODUCTION



IMPLEMENTATION



EXAMPLES





CONCLUSIONS

INTRODUCTION





DNA SEQUENCING



THE PAST

Capillary sequencer.

- 2 Mega (2×10^6) bases a day.
- Read length of about 1000 base pairs




THE PRESENT

SBS (sequencing by synthesis, Novaseq X).

- 8 Terabases (8×10^{12}) a day per sequencer.
- 2.5 bytes a base or 20TB per sequencer per day working space.
- 150 base pair read length
- We have 20 of the previous generation currently in production.



THE FUTURE




<https://www.genome.gov/about-genomics/fact-sheets/DNA-Sequencing-Costs-Data>

COMPUTE CLUSTERS

- A number of very similar configuration LSF clusters of different sizes. Separate clusters are used to ringfence resources.
- All* lustre filesystems are mounted everywhere.
- Lustre filesystems are not generally dedicated to a single group.

- Cancer and somatic mutation
- Main compute cluster
- General purpose and training
- Impute service
- Pathogens
- Sequencing research
- Sequencing Production
- Tree of Life



```
jb23@seqfarm3-head1:~$ df -H /lustre/*
Filesystem      Size  Used Avail Use% Mounted on
172.17.220.18@tcp:172.17.220.19@tcp:/lus17  3.1P  2.2P  886T  71% /lustre/scratch117
172.17.220.28@tcp:172.17.220.29@tcp:/lus18  3.1P  1.8P  1.3P  59% /lustre/scratch118
172.28.64.102@tcp:172.28.64.103@tcp:/lus19  2.5P  2.2P  254T  90% /lustre/scratch119
10.177.252.5@tcp:10.177.252.4@tcp:/lus20   3.1P  1.2P  1.9P  38% /lustre/scratch120
10.160.32.5@tcp:10.160.32.4@tcp:/lus23   5.3P  2.8P  2.5P  53% /lustre/scratch123
10.160.36.4@tcp:10.160.36.5@tcp:/lus24   5.3P  1.8P  3.5P  34% /lustre/scratch124
10.160.40.4@tcp:10.160.40.5@tcp:/lus25   5.3P  229T  5.1P  5% /lustre/scratch125
10.160.42.4@tcp:10.160.42.5@tcp:/lus26   5.3P  327T  5.0P  7% /lustre/scratch126
```


JOB STATS

- https://doc.lustre.org/lustre_manual.xhtml#jobstats
- A method of tagging I/O with additional metadata.
- Introduced in Lustre 2.3 2011
- A 32 character tag.
- The default tag is procname_uid
 - executable_name.uid
- A filesystem or a compute node can be configured to read a environment variable as the tag.



REQUIREMENTS


- User name
- Group name
- LSF project
- Compute cluster
- A pony



IMPLEMENTATION



ARCHITECTURE



RAW DATA

- Raw data is fetched in grey with lctl get_param
- LU-16110, LU-13857 – jobs_stats not valid yaml
- LU-15870 – job_stats is sometimes corrupt
- LU-15826 – job_id not quoted

```
lctl get_param obdfilter.lus24-0ST0000.job_stats
obdfilter.lus24-0ST0000.job_stats=
job_stats:
- job_id:      SA1CUzMGBwYH,analysis-cgp
  snapshot_time: 1667904951
  read_bytes: { samples: 297356696, unit: bytes, min: 4096, max: 16777216, sum: 1762446621732864, sumsq: 9915053801461514240 }
  write_bytes: { samples: 20518685, unit: bytes, min: 1, max: 16777216, sum: 98655040843302, sumsq: 4627393946935432142 }
  read: { samples: 297356696, unit: usecs, min: 1, max: 12945252, sum: 11870681362787, sumsq: 3234937531575968175 }
  write: { samples: 20518685, unit: usecs, min: 1, max: 10280151, sum: 98226898749, sumsq: 50535912682848539 }
  getattr: { samples: 0, unit: usecs, min: 0, max: 0, sum: 0, sumsq: 0 }
  setattr: { samples: 5697, unit: usecs, min: 4, max: 6119237, sum: 151065923, sumsq: 327770191069603 }
  punch: { samples: 1458531, unit: usecs, min: 5, max: 31397444, sum: 15753447595, sumsq: 36605999152526109 }
  sync: { samples: 6517520, unit: usecs, min: 0, max: 7189473, sum: 10441362075, sumsq: 12097924287750305 }
  destroy: { samples: 0, unit: usecs, min: 0, max: 0, sum: 0, sumsq: 0 }
  create: { samples: 0, unit: usecs, min: 0, max: 0, sum: 0, sumsq: 0 }
  statfs: { samples: 0, unit: usecs, min: 0, max: 0, sum: 0, sumsq: 0 }
  get_info: { samples: 0, unit: usecs, min: 0, max: 0, sum: 0, sumsq: 0 }
  set_info: { samples: 0, unit: usecs, min: 0, max: 0, sum: 0, sumsq: 0 }
  quotactl: { samples: 0, unit: usecs, min: 0, max: 0, sum: 0, sumsq: 0 }
  prealloc: { samples: 0, unit: reqs }
```

WEB APP FOR LUSTRE SERVERS

- Minimal server
- Fix yaml errors
- JSON output for each mounted device

```
def do_GET(self):  
    data={}  
    self._set_headers()  
    osts=map( lambda x: x.replace("/proc/fs/lustre/obdfilter/", ""),glob.glob("/proc/fs/lustre/obdfilter/lus*"))  
    mdts=map( lambda x: x.replace("/proc/fs/lustre/mdt/", ""),glob.glob("/proc/fs/lustre/mdt/lus*"))  
    for ost in osts:  
        data[ost]=run_lctl('lctl get_param obdfilter.{0}.job_stats'.format(ost).split())  
    for mdt in mdts:  
        data[mdt]=run_lctl('lctl get_param mdt.{0}.job_stats'.format(mdt).split())  
    self.wfile.write(json.dumps(data))
```

WEB APP FOR LUSTRE SERVERS

```
fix_jobname = re.compile(r'^- job_id:(\s*)(\S(?:.*\S)?)\s*\$')

def process_line(line):
    # Quote the job_id
    line = fix_jobname.sub(r'- job_id:\1"\2"',line)
    # Ensure valid yaml with a space
    return line.replace("max:", "max: ").replace("min:", "min: ").replace("sum:", "sum: ").replace("samples:", "samples: ")
.replace("sumsq:", "sumsq: ")

def run_lctl(cmd):
    p = Popen(cmd, stdout=PIPE)
    # skip the first two headers lines
    # https://stackoverflow.com/questions/15571137/yaml-scanner-scannererror-while-scanning-a-directive
    results=[]
    data_stream = "\n".join(map(process_line,p.communicate()[0].split("\n")[2:]))
    return yaml.load(data_stream)
```

SAMPLE OUTPUT

```
curl -s http://lus23-oss1-mgmt.internal.sanger.ac.uk:8081 | jq .
{
  "lus23-OST0001": [
    {
      "write_bytes": {
        "max": 16777216,
        "sum": 1275960319678,
        "samples": 413177,
        "unit": "bytes",
        "min": 4
      },
      "job_id": "SA1TwDkw0zA7,team135",
      "setattr": {
        "samples": 7,
        "unit": "reqs"
      },
      ...
      "read_bytes": {
        "max": 16777216,
        "sum": 51588627873792,
        "samples": 22190157,
        "unit": "bytes",
        "min": 4096
      },
      "snapshot_time": 1667913397
    }
  ]
}
```




IS 32 BYTES ENOUGH?

SA1|UzMGBwYH,analysis-cgp

SA1 The version string, used to confirm format
C The compute farm (CASM)
UzMGBwYH base64 encoded (uid,pgid,gid)
, End of fixed data
analysis-cgp LSF project

```
echo UzMGBwYH | base64 --decode | od -x
0000000 3353 0706 0706
jb23@farm5-head1:~$ id cgppipe
uid=13139(cgppipe) gid=1798(cancer-pipeline) groups=1798(cancer-
pipeline),15495(metanorm),1300(cancer_external),1722(musperm),1116(team113),80
8(cancer),9215(team78),1534(cgppc),1410(cgp_transfer),1554(grp14mg)
jb23@farm5-head1:~$ bc
bc 1.07.1
(3*16+3)*256+5*16+3
13139
7*256+6
1798
```



ESUB - JOB PREFLIGHT SCRIPT

```
if var_has_contents("LSB_SUB MODIFY_ENVFILE"):
    with open(os.environ["LSB_SUB MODIFY_ENVFILE"], "a") as modfile:
        header = "SA1"
        # Write a char which is the cluster encoded, U is unknown.
        if var_has_contents("CLUSTER"):
            cluster = encode_farm.get(os.environ["CLUSTER"].splitlines()[0] "U")
        else:
            cluster = "U"
        uids = base64.b64encode(
            struct.pack(
                b"hhh",
                os.getuid(),
                pwd.getpwuid(os.getuid()).pw_gid,
                os.getgid()
            )
        ).decode().rstrip("=")
        if var_has_contents("LSB_PROJECT_NAME"):
            if os.environ["LSB_PROJECT_NAME"] == "default":
                project = os.environ["LSB_DEFAULTGROUP"]
            else:
                project = os.environ["LSB_PROJECT_NAME"]
        else:
            if var_has_contents("LSB_DEFAULTGROUP"):
                project = os.environ["LSB_DEFAULTGROUP"]
            else:
                project = ""
        modfile.write(SANGER_JOB_ID="" + header + cluster + uids + ',' + project + "\n")
        modfile.write(SINGULARITYENV_SANGER_JOB_ID="" + header + cluster + uids + ',' + project + "\n")

if [ -n "${STAT_TYPE}" ] ; then
    lctl set_param jobid_var=${STAT_TYPE} >/dev/null
fi
```




WEB APP FOR LUSTRE SERVERS

```
def decode_jobname(name, store):
    store['project'] = "unknown"
    store['uid'] = "unknown"
    store['pgid'] = "unknown"
    store['gid'] = "unknown"
    store['farm'] = "unknown"
    decode_farm = {
        "F": "farm5",
        "C": "casm3",
        "G": "gen3",
        "I": "impute",
        "i": "isg-test-cluster",
        "P": "pcs6",
        "S": "seq4",
        "s": "seqfarm3",
        "T": "tol"
    }
    if name.startswith("SA1"):
        store['farm'] = decode_farm.get(name[3], "unknown")
        parse = re.search("^(SA1.([^\n]*),(.*)$)", name)
        if parse is not None:
            try:
                (uid, pgid, gid) = struct.unpack('hhh', base64.b64decode(parse.group(1) + "==="))
                store['project'] = parse.group(2) or ""
                store['uid'] = str(uid)
                store['pgid'] = str(pgid)
                store['gid'] = str(gid)
            except:
                pass
```



REAL WORLD

```
curl -s http://lus23-oss1-mgmt.internal.sanger.ac.uk:8080 | jq .
{
  "lus23-OST0001": [
    {
      "write_bytes": {
        "max": 16777216,
        "sum": 1277816778962,
        "samples": 413738,
        "unit": "bytes",
        "min": 4
      },
      "pgid": "15152",
      "job_id": "SA1TwDkw0zA7,team135",
      "uid": "14784",
      ...
      "project": "team135",
      "gid": "15152",
      "farm": "tol",
      "read_bytes": {
        "max": 16777216,
        "sum": 51752874618880,
        "samples": 22381014,
        "unit": "bytes",
        "min": 4096
      }
    }
  ]
}
```



CONFUSION

- We see corruption in the job id, LU-15870, LU-16251
- We see kworker recorded as the executable with root.
- We see thread-pool with no uid.

```
- job_id:      kworker/u518:6eamtrynka
snapshot_time: 1668009798
read_bytes:   { samples:          1, unit: bytes, min:  663552, max:  663552,
sum:          663552 }
write_bytes:  { samples:          0, unit: bytes, min:      0, max:      0,
sum:          0 }

- job_id:      kworker/u130:0
snapshot_time: 1668009796
read_bytes:   { samples:         21459, unit: bytes, min:  4096, max: 2297856,
sum:          160104448 }
write_bytes:  { samples:          0, unit: bytes, min:      0, max:      0,
sum:          0 }

- job_id:      thread-pool-14113912
snapshot_time: 1668009775
read_bytes:   { samples:          2, unit: bytes, min: 36864, max: 983040,
sum:          1019904 }
write_bytes:  { samples:          0, unit: bytes, min:      0, max:      0,
sum:          0 }

[root@lus23-oss1 ~]# lctl get_param obdfilter.lus23-OST0000.job_stats | grep job_id
| grep thread-pool | wc -l
126
[root@lus23-oss1 ~]# lctl get_param obdfilter.lus23-OST0000.job_stats | grep job_id
| grep kworker | wc -l
155
[root@lus23-oss1 ~]# lctl get_param obdfilter.lus23-OST0000.job_stats | grep job_id
| wc -l
330
```



PUSHING METRICS TO GRAPHITE

- Fetches data from all OSS's and MDS's
- Adds human readable versions of UID and GUD

```
def uid_to_name(uid: str) -> str:  
    try:  
        return pwd.getpwuid(int(uid)).pw_name  
    except (KeyError, ValueError):  
        return "unknown"
```

```
def gid_to_name(gid: str) -> str:  
    try:  
        return grp.getgrgid(gid)[0]  
    except (KeyError, ValueError):  
        return "unknown"
```

- Generates multiple views
 - Per Farm
 - Per LSF project
 - Per User
 - Per Group
- Send data to Graphite



SECURE LUSTRE

- A method of delivering lustre filesystems to OpenStack projects using
 - Sub directory mounts
 - And user id mappings
- <https://hpc-news.sanger.ac.uk/2021/05/20/update-on-secure-lustre-lustre-users-group-may-2020/>
- Jobid is inserted on the client, the uid inserted may and often does not exist on the outside
- The nodemap the request is received on is unavailable in the jobstats data
- Our current aggregation program appears to drop this traffic.



EXAMPLES





filesystem

lus23

type

Enter variable value

byfarm

bygroup

bylsf_project

byusername

NB occasional u

Compare with overall I/O showr



The red line indicates 90% of te

ata "spikes" are due to a Lustre server bug and should be disregarded.


internal.sanger.ac.uk/d/000000203/lustre-stats?orgId=1



ty. The yellow line indicates 85%. Around these levels, the filesystem will feel slow to interactive users.



▼ Read/Write



▼ Stats







CONCLUSIONS



CONCLUSIONS

- Small change on Lustre server
- Not perfect
- Still useful
 - Want fair share for I/O
 - Is NRS the answer to that question
- Released under Apache 2 license





ACKNOWLEDGMENTS

- ISG (Photo is missing Anna, Bruno, Fabio and Sai) (And includes Brett who is no longer with us).

THANK YOU



@wellcomegenomecampus



@wellcomegenome



Wellcome Genome Campus



Wellcome Genome Campus